

Order-Dependent Bias in Sequential Explicit Euler Physiological Simulation: Evidence from a Multi-Organ Canine Cardiovascular Platform

Authors: [Author list]Corresponding: [Email]Keywords: sequential Euler coupling, baroreflex simulation, order-dependent bias, multi-organ physiological modeling, implicit solver recommendation

Abstract

English

The baroreflex maintains arterial pressure through coordinated heart and neurohumoral responses, but the impact of module ordering in sequential physiological simulation has not been systematically characterized. Implementing a baroreflex model with two explicit Euler update sequences revealed order-dependent bias that is both time-step invariant and scenario-dependent. At baseline conditions (cardiac output = 5.0 L/min, systemic vascular resistance = 17.5 mmHg/L/min), sequential heart→neuro updates introduced a 44.7 mmHg mean arterial pressure (MAP) error, while the neuro→heart ordering produced near-zero error. However, during 400 mL hemorrhage recovery at $t = 30$ s, the accuracy rankings reversed: heart→neuro matched the implicit Radau solver reference with $\Delta = 0.1$ mmHg, whereas neuro→heart underestimated MAP by 9.4 mmHg. Error analysis confirmed first-order convergence breakdown — refining the time step from 0.1 s to 1×10^{-9} s did not reduce the bias. A unified Euler formulation coupling all organ modules simultaneously yielded $\text{RMSE} < 0.2$ mmHg regardless of initialization order. These results demonstrate that Gauss-Seidel information lag in sequential coupling creates order-dependent bias that renders any fixed module ordering unreliable across physiological states. A unified state-vector formulation — in which all organ modules evaluate derivatives against the identical state vector simultaneously — is required for reproducible baroreflex simulation; implicit solvers or small-step explicit integration applied to the unified RHS achieve $\text{RMSE} < 0.2$ mmHg.

Words: 247

中文

压力感受性反射通过协调心脏和神经体液响应维持动脉血压，但顺序生理学模拟中模块排序的影响尚未得到系统表征。在压力感受性反射模型中采用两种显式欧拉更新序列进行实现，发现排序相关偏差具有时间步长无关性和场景依赖性。在基线条件下（心输出量 = 5.0 L/min，系统血管阻力 = 17.5 mmHg/L/min），顺序的心脏→神经更新引入 44.7 mmHg 平均动脉压（MAP）误差，而神经→心脏排序产生的误差接近零。然而，在 400 mL 出血恢复后 $t = 30$ s 时，准确度排序发生逆转：心脏→神经更新与隐式 Radau 求解器参考值匹配， $\Delta = 0.1$ mmHg，而神经→心脏排序低估 MAP 达 9.4 mmHg。误差分析证实一阶收敛性失效——将时间步长从 0.1 s 细化至 1×10^{-9} s 未能减小偏差。统一欧拉公式同时耦合所有器官模块，无论初始化顺序如何均达到

RMSE < 0.2 mmHg。这些结果表明，高斯-塞德尔信息滞后产生不对称的执行器-传感器延迟，使任何固定模块排序在生理状态间均不可靠。隐式求解器同时求解所有器官方程，方可实现可重复的压力感受性反射模拟。

字数：241

1. Introduction

Multi-organ physiological simulation is increasingly used for veterinary and medical education, clinical decision training, and research mechanistics. Computational physiology platforms must balance biophysical fidelity with numerical tractability, and many use explicit Euler with sequential module coupling for simplicity and speed. The baroreflex — a classic closed-loop system linking arterial pressure sensing to heart rate and vascular tone — is a particularly relevant test case because it is a tight, two-module feedback loop with asymmetric timescales: the cardiac actuator responds within one to three seconds, while the neurohumoral sensor operates with a 5–10 second time constant. This asymmetry is precisely what makes the baroreflex loop vulnerable to sequential coupling bias.

Despite the widespread use of sequential explicit Euler in physiological modeling, no prior work has systematically quantified how module update ordering affects the equilibrium and transient behavior of multi-organ cardiovascular simulation. A 2024 study by T□a□ka et al. published in PLOS Computational Biology noted that explicit Euler methods are "numerically unstable" for baroreflex simulation and recommended implicit solvers such as Tsit5() and delay differential equations (DDE), yet the field lacks a systematic characterization of the order-dependent bias inherent in sequential coupling. This gap is consequential: for an educational or diagnostic platform, a ± 45 mmHg mean arterial pressure error at baseline represents the difference between a clinically plausible and a severely implausible vital sign reading.

In this paper we use Virtual Vet — an 11-organ canine cardiovascular simulation platform — as our test environment. We systematically compare three coupling strategies (unified Euler, sequential Euler with heart→neuro ordering, sequential Euler with neuro→heart ordering) against an implicit Radau solver reference. We present two physiological scenarios: a baseline equilibrium (no perturbation) and a 400 mL hemorrhage transient (acute blood loss beginning at $t = 5$ s). Our central finding is that sequential coupling introduces order-dependent bias whose magnitude and direction vary unpredictably with physiological condition — no fixed module ordering is safe across all scenarios.

2. Methods

2.1 Virtual Vet Platform Architecture

Virtual Vet is an 11-organ canine cardiovascular simulation built in Python 3. The platform integrates organ modules for heart, lung, kidney, gut, liver, endocrine, neuro (baroreflex), immune, coagulation, lymphatic, and fluid. All modules share a blood compartment, and all physiological parameter modifications are applied exclusively through a FactorCommand interface that prevents direct attribute mutation. The default integration time step is $dt = 0.01\text{ s}$; all module state variables are stored in a unified state vector for implicit solver access.

Table 1: Virtual Vet 11-organ module inventory

Module	Key State Variables	Physiological Role
heart	HR, SV, SVR, circulating_volume_ml	Cardiac output, MAP computation
lung	PaO_2, PaCO_2, pH	Gas exchange, ventilatory drive
kidney	GFR, urine_flow, plasma_ADH	Fluid balance, RAAS
neuro	sympathetic_tone, parasympathetic_tone	Baroreflex integration
fluid	plasma_vol, interstitial_vol	3-compartment volume dynamics
gut	absorption_rate	Nutrient absorption
liver	metabolic_rate	Glycogen stores
endocrine	cortisol, insulin	Hormone axes
immune	TNF_alpha, IL6	Inflammatory response
coagulation	fibrinogen, platelets	Clotting cascade
lymphatic	lymph_flow	Interstitial drainage

2.2 Heart Module and Baroreflex Loop A

The heart module maintains four state variables: heart rate (HR, beats/min), stroke volume (SV, mL/beat), systemic vascular resistance (SVR, mmHg/L/min), and circulating blood volume (BV, mL). Mean arterial pressure is computed from the Frank-Starling and Windkessel relationships:

$$CO = HR \times SV / 1000 \text{ (cardiac output, L/min, SV in mL)}$$
$$MAP = MAP_base + (CO / 60) \times SVR \text{ (1)}$$
$$SV = base_SV \times f(BV / total_BV) \text{ (2)}$$

The baroreflex loop A operates within the heart module itself: MAP error drives changes in HR and SVR through sympathetic and parasympathetic channels with separate gain and time-constant parameters (heart.py derivatives, lines 163–187).

2.3 Neuro Module and Baroreflex Loop B

The neuro module maintains sympathetic_tone and parasympathetic_tone as state variables (0–1 scale). It receives MAP, PaO_2, PaCO_2, and pH as inputs and computes target sympathetic drive

through pain, seizure, and chemoreceptor pathways (neuro.py derivatives, lines 121–131). Its compute() method emits FactorCommand operations targeting heart.heart_rate and heart.SVR, applied to the heart module before the next heart.compute() call in the sequential path.

Architecture: Two independent baroreflex mechanisms operate simultaneously. Loop A (intra-heart, heart.py derivatives lines 167–193): MAP error drives HR/SVR via the heart's own sympathetic and parasympathetic state variables (heart.py lines 83–84), which are packed into the unified state vector. Loop B (inter-module, neuro.compute() → FactorCommands → heart): the neuro module's computed sympathetic tone writes to heart.heart_rate and heart.SVR through a separate pathway. The sequential coupling bias arises from the ordering of when loop B's FactorCommands are applied relative to when loop A computes its response from MAP error.

2.4 Three Coupling Strategies Compared (Plus One Reference)

We compare three coupling strategies — unified Euler, sequential Euler with heart→neuro ordering, and sequential Euler with neuro→heart ordering — plus one implicit reference method as the gold-standard comparator. The three strategies differ only in update order; the reference method differs in solver type.

Method 1 — Unified State-Vector + Explicit Solver: All 11 organ modules evaluate derivatives against the same state vector via VirtualCreature._unified_rhs(t, y), which is passed to scipy.integrate.solve_ivp (method='RK45', rtol=1e-6). Because all modules see the identical state simultaneously, there is no Gauss-Seidel information lag regardless of solver choice. We confirmed that RK45 and fixed-step forward Euler (dt=0.01 s) produce equivalent accuracy (RMSE < 0.2 mmHg in both cases), establishing that the method's accuracy comes from simultaneous coupling rather than solver order.

Method 2 — Sequential Euler (heart→neuro): heart.compute() is called first, updating HR, SV, and SVR based on the current neuro state from the previous time step. Then neuro.compute() is called, seeing the updated cardiovascular state from the current time step. This is a Gauss-Seidel iteration where the baroreceptor (neuro) leads the actuator (heart).

Method 3 — Sequential Euler (neuro→heart): The update order is reversed. neuro.compute() first sees the cardiovascular state from the previous time step (actuator lag), then heart.compute() sees the updated neuro state from the current time step.

Method 4 — Radau Reference: The unified RHS is passed to scipy.integrate.solve_ivp with method='Radau', rtol = 1e-10, atol = 1e-12, as the gold-standard reference. Radau is a fully implicit 5th-order Runge-Kutta method that solves all ODEs simultaneously.

2.5 Hemorrhage Model

Blood loss is modeled as a sigmoid forcing term applied to the blood volume derivative:

```
blood_loss_rate_ml_s = k x sigmoid_on(t; t_onset, width) x sigmoid_off(t; t_onset+duration, width)
```

For the 400 mL hemorrhage experiment: t_onset = 5 s, total_ml = 400 mL, duration = 300 s, width = 6 s, k = 35 mL/s. This produces a gradual blood volume reduction over approximately 20 seconds beginning at t = 5 s.

2.6 Convergence Study Design

Pure Euler was evaluated at $dt = [0.1, 0.05, 0.025, 0.01, 0.005, 0.002, 0.001, 0.0005, 0.0002, 0.0001]$ s on a 60 s observation window. Reference was BDF with $rtol = 1e-6$, $max_step = 0.5$ s. RMSE (MAP) was computed at every 0.5 s time point. Convergence order was estimated from $\log(RMSE)$ vs $\log(dt)$ regression.

To confirm that the sequential Euler bias is $O(1)$ rather than a numerical truncation error, a dedicated experiment was conducted at $dt = 1 \times 10^{-9}$ s under the heart→neuro ordering (60 s simulation, 60 million steps). If the bias were a truncation error, halving dt by a factor of 10^7 should reduce it proportionally; if it is $O(1)$, the bias should remain ± 44.7 mmHg. The result confirmed $O(1)$: MAP at $t = 60$ s was 144.7 mmHg (Figure 5), identical to the $dt = 0.01$ s result within machine precision, establishing that the bias does not vanish as $dt \rightarrow 0$.

3. Results

3.1 Baseline Equilibrium: Sequential Coupling Introduces $O(1)$ Bias

At baseline (no hemorrhage, 20 kg canine, $dt = 0.01$ s), sequential Euler coupling with the heart→neuro ordering produced a sustained MAP of 144.7 mmHg — a +44.7 mmHg error relative to the 100.0 mmHg reference (Table 2). The neuro→heart ordering produced MAP = 100.0 mmHg, matching the reference. This 44.7 mmHg bias is clinically implausible (severe hypertension) and is not reduced by refining the time step: testing across $dt = 0.1$ s to $dt = 1 \times 10^{-9}$ s confirmed that the bias is $O(1)$ and time-step invariant.

Table 2: Baseline MAP at $t = 60$ s, Sequential Euler Order Comparison

Ordering	MAP (mmHg)	Delta vs Reference	Bias Direction
heart→neuro	144.7	+44.7	Over-compensation
neuro→heart	100.0	± 0.0	Accurate
Radau reference	100.0	—	—

The pure unified Euler method (Method 1) achieved $RMSE < 0.2$ mmHg regardless of which module conceptually "leads," confirming that the bias is inherent to sequential coupling and not to the physiological model itself. Convergence analysis confirmed first-order convergence for pure Euler: $RMSE \propto dt$, consistent with standard explicit Euler theory (Figure 2).

Table S1 (Supplementary): Pure Euler RMSE vs dt — 60 s Observation, 400 mL Hemorrhage

dt (s)	RMSE MAP (mmHg)	Convergence Order
0.1	1.847	—
0.05	0.938	0.98
0.025	0.473	0.99

dt (s)	RMSE MAP (mmHg)	Convergence Order
0.01	0.191	0.99
0.005	0.096	1.00
0.001	0.019	1.01
0.0001	0.002	1.00

The swap experiment provides causal proof that the bias originates from the sequential coupling architecture. When module order was reversed, the bias direction also reversed: DeltaMAP switched from +44.7 mmHg (heart→neuro) to +/-0.0 mmHg (neuro→heart), as shown in Figure 3. This reversal cannot be explained by parameter variation — it is a direct consequence of which module sees updated state and which sees the previous time step's state.

3.2 Hemorrhage Transient: The Central Paradox — Ordering Correctness Reverses

Under 400 mL hemorrhage (onset $t = 5$ s), the baseline accuracy rankings reversed (Table 3 and Figure 4). Both orderings reached nearly identical MAP_nadir during the acute phase (Delta = 0.2 mmHg at MAP_min ~ 89 mmHg), but the recovery trajectory diverged sharply: at $t = 30$ s, heart→neuro recovered to 98.4 mmHg (Delta = +0.1 mmHg vs reference) while neuro→heart remained suppressed at 89.0 mmHg (Delta = -9.4 mmHg vs reference). By $t = 120$ s, both orderings had reconverged to approximately 97.4 mmHg.

Table 3: Hemorrhage Transient Metrics, 400 mL Blood Loss

Metric	heart→neuro	neuro→heart	Reference	Error vs Reference	Delta (h→n vs n→h)
MAP_min (mmHg)	89.2	89.0	~ 89.2	+0.0 / -0.2	0.2
t_MAP_min (s)	24.9	28.0	~ 25.0	—	3.1
MAP @ 30s (mmHg)	98.4	89.0	~ 98.4	+0.1 / -9.4	9.4
MAP @ 120s (mmHg)	97.4	97.4	~ 97.4	+0.0 / +0.0	0.0

The "Error vs Reference" column shows (heart→neuro error, neuro→heart error) in mmHg relative to the reference. Reference values are from the Radau implicit solver (method='Radau', rtol=1e-10, atol=1e-12) cross-validated with BDF (max difference < 0.01 mmHg across all time points).

The paradox: The ordering that was accurate at baseline (neuro→heart, +/-0.0 mmHg error) became the inaccurate one during hemorrhage recovery (-9.4 mmHg at $t = 30$ s). The ordering that was catastrophically wrong at baseline (heart→neuro, +44.7 mmHg error) became the accurate one during recovery (+0.1 mmHg at $t = 30$ s). This is the central finding of this paper.

3.3 Order-Independence of Pure Unified Euler

Pure unified Euler with the same $dt = 0.01$ s achieved MAP at $t = 30$ s within 0.2 mmHg of the Radau reference regardless of initialization order. This confirms that the $O(1)$ bias seen in sequential Euler is not a property of the physiological model — it is entirely a property of the sequential coupling architecture. When all organ modules evaluate derivatives against the same state vector simultaneously, no module "leads" or "lags" and the baroreflex feedback loop functions correctly.

4. Discussion

4.1 Mechanism: Dual Baroreflex Loop Coupling and the SVR Accumulation Feedback

The architecture clarification in §2.3 changes how the mechanism must be understood. Two independent baroreflex mechanisms operate simultaneously: loop A (intra-heart, `heart.py` derivatives lines 167–193) computes HR and SVR targets from MAP error using the heart's own sympathetic/parasympathetic state; loop B (inter-module, `neuro.compute()` → `FactorCommands`) emits HR additive and SVR multiplicative commands targeting `heart.heart_rate` and `heart.SVR`. In the sequential Euler path, `FactorCommands` from loop B are applied to the heart module before the next `heart.compute()` call.

Under `heart→neuro` ordering, `heart.compute()` runs first: loop A computes HR/SVR targets from the current MAP and writes them into `self.heart_rate` and `self.SVR`. Critically, loop A's SVR target `SVR_increase = 1.0 + 2.0 * self.sympathetic * max(0.0, error)` (`heart.py` line 183) starts from the baseline SVR. Then `neuro.compute()` runs second, sees the updated MAP (now slightly elevated by the SVR step), and computes loop B's `FactorCommands` — but loop B's SVR multiplicative command is applied after loop A has already committed to its SVR target for this step. Loop A does not react to loop B's SVR modification within the same step — it reacts only at the following step, compounding the elevation. Over 60 steps at $dt = 0.01$ s, the unopposed SVR elevation accumulates to $\sim +60\%$ above baseline, driving MAP from 100 toward 144.7 mmHg (Figure 3B). This is the actuator-leading configuration: heart's SVR adjustment acts first and dominates, while neuro's corrective signal arrives one step late and is over-written at each iteration.

Under `neuro→heart` ordering, `neuro.compute()` runs first using the stale MAP from the previous step. At baseline (MAP ~ 100 , error ~ 0), loop B's `FactorCommands` are near-zero — the stale MAP produces negligible sympathetic drive. Then `heart.compute()` runs and sees the correct current MAP (100 mmHg), so loop A's SVR target is also computed near baseline (error ~ 0). Both loops agree and MAP stays at 100 mmHg. This is the sensor-leading configuration: the neuro sensor's stale MAP prevents any spurious SVR command from being issued, and heart's loop A responds to the accurate current MAP. The bias reversal during hemorrhage follows from the same mechanism in reverse (see below).

The asymmetry reverses during hemorrhage because the dominant physics changes: blood volume loss drives MAP downward, and error becomes large and positive. Under `heart→neuro`, both loops agree on the need for SVR compensation — heart's loop A elevates SVR from baseline, and when

neuro then computes its FactorCommands from the same updated MAP, the two signals are coherent and produce the correct compensatory response (MAP recovers to 98.4 mmHg at $t = 30$ s, $\Delta = +0.1$ mmHg). Under neuro→heart, however, neuro computes first from the stale pre-hemorrhage MAP (still near 100 mmHg before the volume effect propagates), so loop B's SVR command is issued based on a MAP that does not yet reflect the blood loss. Then heart's loop A runs second: its SVR target is computed using the already-elevated SVR from loop B's command as the base, and the compounded elevation causes an overshoot in the SVR trajectory, keeping MAP suppressed at 89.0 mmHg ($\square 9.4$ mmHg error). The reversal occurs because the sensor-lag configuration (neuro→heart), which was benign at baseline (error ~ 0 , stale MAP produced near-zero commands), becomes harmful when error is large (stale MAP produces an SVR command that is too small or too large depending on the hemodynamic context). In both orderings the same mechanism operates — the ordering of when loop A computes its SVR target relative to when loop B's SVR modification is applied — but the net effect on the MAP trajectory reverses depending on whether MAP is rising toward an erroneously elevated equilibrium or falling toward a physiologically depressed nadir.

The control theory framing — actuator lag vs sensor lag — is qualitatively suggestive but imprecise for this system. A rigorous treatment would require linearizing the coupled SVR dynamics around the operating point; the qualitative framing is kept here for pedagogical clarity, pending that formal analysis.

4.2 Comparison with Prior Art

Taika et al. 2024 (PLoS Computational Biology) performed the first global sensitivity analysis of closed-loop baroreflex regulation in pulsatile cardiovascular models, finding that baroreflex parameters substantially influence cardiac output under closed-loop operation. Their work underscores that the baroreflex feedback architecture critically determines simulation output — consistent with our finding that coupling strategy (order, explicit vs. implicit) determines MAP trajectory. Ottesen and Olufsen's 2004 textbook multi-organ models recommended implicit solvers for closed-loop cardiovascular simulation; our work provides quantitative evidence for why this recommendation is necessary and demonstrates that the consequence of ignoring it is not just numerical instability but systematic, order-dependent distortion of physiological output. Ursino's 1998 baroreflex model used implicit Gear-style solvers, representing the historical precedent that supports our recommendation.

This order-dependent failure parallels the well-known added-mass instability in partitioned fluid-structure interaction (Causin, Gerbeau & Nobile 2005; Förster, Wall & Ramm 2007), where staggered coupling diverges for certain density ratios regardless of time-step refinement. In both cases, sequential coupling can produce order-dependent bias that is unrecoverable by time-step refinement — a unified formulation that evaluates both sides against a shared state is required to recover the correct solution. Guyton et al. (1972) established the classical whole-body cardiovascular model upon which many modern educational simulators are built; the numerical challenges documented here apply directly to any such multi-module platform using sequential coupling. Strang (1968) formalized operator splitting as a general framework, noting that symmetric (Strang) splitting is required for second-order accuracy — a principle

that foreshadowed the ordering sensitivity we document here. Shi, Udelson et al. (2011) reviewed numerical methods for OD cardiovascular modeling, noting that explicit schemes require impractically small time steps for stable coupling of stiff closed-loop models; our results extend this observation to show that even simultaneous explicit evaluation of all modules (unified Euler) recovers correct behavior where sequential evaluation does not.

4.3 Implications for Educational and Clinical Simulation

Virtual Vet and similar platforms used for veterinary or medical education must disclose the coupling strategy and module ordering as fundamental methodological facts, not implementation details. A ± 45 mmHg MAP error at baseline would produce incorrect vital sign interpretations in a clinical training scenario. Developers of multi-organ physiological simulations should default to implicit solvers (Radau or BDF) with a unified right-hand-side formulation, or at minimum conduct and report a sensitivity analysis across at least two module orderings.

4.4 Recommendation

Primary recommendation: Unified state-vector formulation (all modules read from and write to the same state vector simultaneously) with either implicit or small-step explicit integration. This eliminates sequential coupling bias entirely; RMSE < 0.2 mmHg confirmed across all tested scenarios regardless of the underlying ODE solver. For the Virtual Vet platform, `VirtualCreature._unified_rhs(t, y)` passed to `scipy.integrate.solve_ivp` with `method='Radau'` provides the most robust solution.

Secondary recommendation: If only sequential explicit Euler is feasible, run the same simulation with two orderings (heart $\rightarrow$ neuro AND neuro $\rightarrow$ heart) and report the maximum difference across orderings as an uncertainty bound attached to every MAP time series.

Minimum disclosure: Any simulation using sequential coupling must report which module is computed first and provide a sensitivity analysis across at least two orderings.

5. Conclusion

Sequential explicit Euler coupling in a multi-organ baroreflex simulation introduces order-dependent bias that is $O(1)$ (time-step invariant) and causally reverses with module order swap. The bias magnitude and direction vary with physiological condition: at baseline, heart $\rightarrow$ neuro produces a $+44.7$ mmHg MAP error while neuro $\rightarrow$ heart is accurate; during hemorrhage recovery at $t = 30$ s, the rankings reverse and heart $\rightarrow$ neuro becomes accurate while neuro $\rightarrow$ heart underestimates MAP by 9.4 mmHg. A pure unified Euler formulation achieves RMSE < 0.2 mmHg regardless of initialization order.

No module ordering is universally safe across all physiological scenarios. The physiological simulation community should adopt a unified state-vector formulation -- in which all organ modules evaluate derivatives against the identical state vector simultaneously -- as the standard for multi-organ baroreflex simulation, with explicit documentation of which solver is applied to the unified RHS. Any remaining use of sequential coupling must be accompanied by

systematic ordering sensitivity analysis. This finding extends beyond the baroreflex: any tight feedback loop in a multi-organ simulation -- renal RAAS, pulmonary gas exchange, endocrine axes -- is potentially subject to sequential coupling bias. Future work should map the bias landscape across the full 11-organ module graph to identify which coupling topologies are robust and which require unified state-vector formulation.

References

Supplementary Materials

Figure S1: Multi-Organ Platform Architecture

Schematic of the Virtual Vet 11-organ platform showing the dual baroreflex loop architecture: loop A (intra-heart, heart.py derivatives) and loop B (inter-module, neuro.compute() -> FactorCommands -> heart). The sequential coupling ordering (heart->neuro vs neuro->heart) determines when loop B's FactorCommands are applied relative to loop A's SVR target computation.

Figure S2: Pure Euler Convergence Analysis

(A) RMSE(MAP) vs dt on semi-log axes for pure Euler at dt = 0.1 to dt = 1x10⁻⁴ s, 60 s observation window. First-order convergence is confirmed: RMSE ∝ dt with slope ~ 1.0. (B) MAP time series showing plateau elimination: pure Euler at dt=0.01 s converges to the same steady-state as finer dt values, unlike sequential Euler which shows a divergent plateau.

Figure S3: Baseline Swap Experiment -- Time Series and Equilibrium Divergence

(A) MAP time series under heart->neuro and neuro->heart orderings at baseline, dt = 0.01 s. neuro->heart maintains MAP = 100.0 mmHg throughout; heart->neuro diverges to 144.7 mmHg by t = 40 s. (B) Equilibrium MAP after 60 s for each ordering, confirming the +/-44.7 mmHg asymmetry.

Figure S4: 400 mL Hemorrhage Transient -- Order Accuracy Reversal

MAP (top) and blood volume (bottom) time series for 400 mL hemorrhage beginning at t = 5 s. Both orderings track similarly during the acute phase (t = 5-25 s) but diverge during recovery (t = 25-40 s): heart->neuro recovers to ~98 mmHg by t = 30 s (accurate) while neuro->heart remains suppressed at ~89 mmHg (□9.4 mmHg error).

Figure S5: Time-Step Invariance of Sequential Euler Bias

MAP at t = 60 s under heart->neuro ordering across seven orders of magnitude in dt (0.1 s to 1x10⁻⁹ s). The bias is constant at +44.7 mmHg across all tested time steps, confirming 0(1) (time-step invariant) bias rather than a numerical truncation error. Each point represents a separate simulation; the flat line at 144.7 mmHg spans dt = 10⁰ to dt = 10⁻⁹.

Table S1: Pure Euler Convergence Data

See Table S1 above in Section 3.1

Table S2: Module Ordering Sensitivity at Baseline

Full time series of MAP under heart→neuro and neuro→heart orderings, 0–60 s, dt = 0.01 s

Code Availability

The Virtual Vet platform source code is available at: [repository URL]. Experiment scripts for the convergence study, order swap experiment, and hemorrhage transient are located in the experiments/ directory.

Corresponding author: [Name], [Institution], [email]Competing interests: None
declared.Funding: [funding source]